

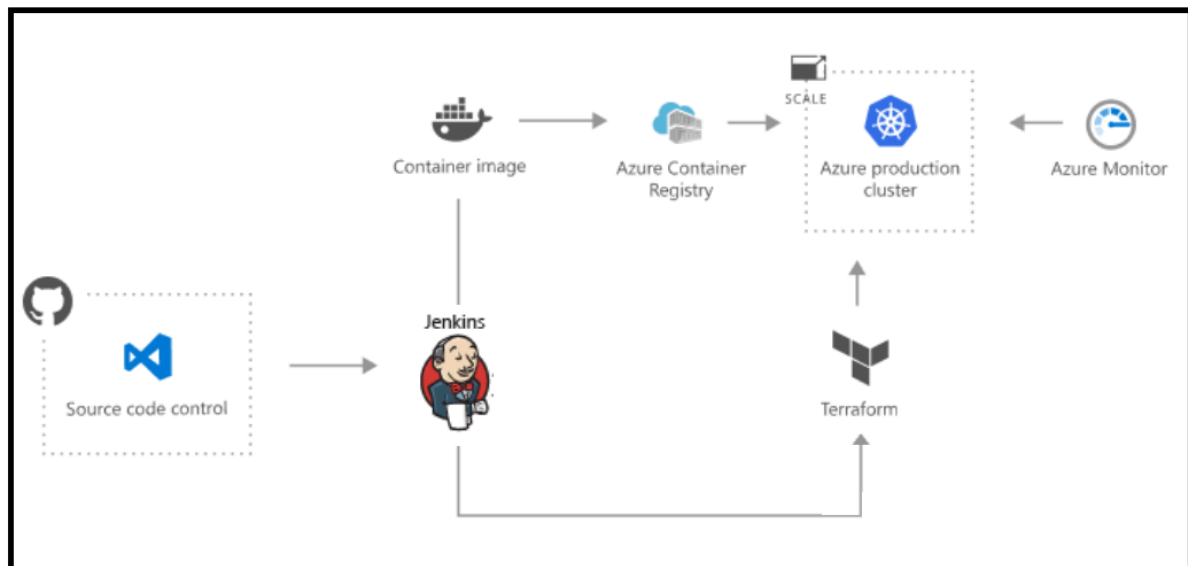
High-level architecture solution and approach to the problem.

Technical Stack Used :

Azure [ACR , AKE , BITNAMI Jenkins Cluster, Azure Cosmo DB] , Terraform , Jenkins

Application Code : Python , Flask (framework) , Beautiful Soup (pkg), Selenium

Note: *Alternatively we can use Azure DevOps instead of Jenkins and Azure Data Factory , to strip the data from the provided datasource [data] . This will simplify the solution.*

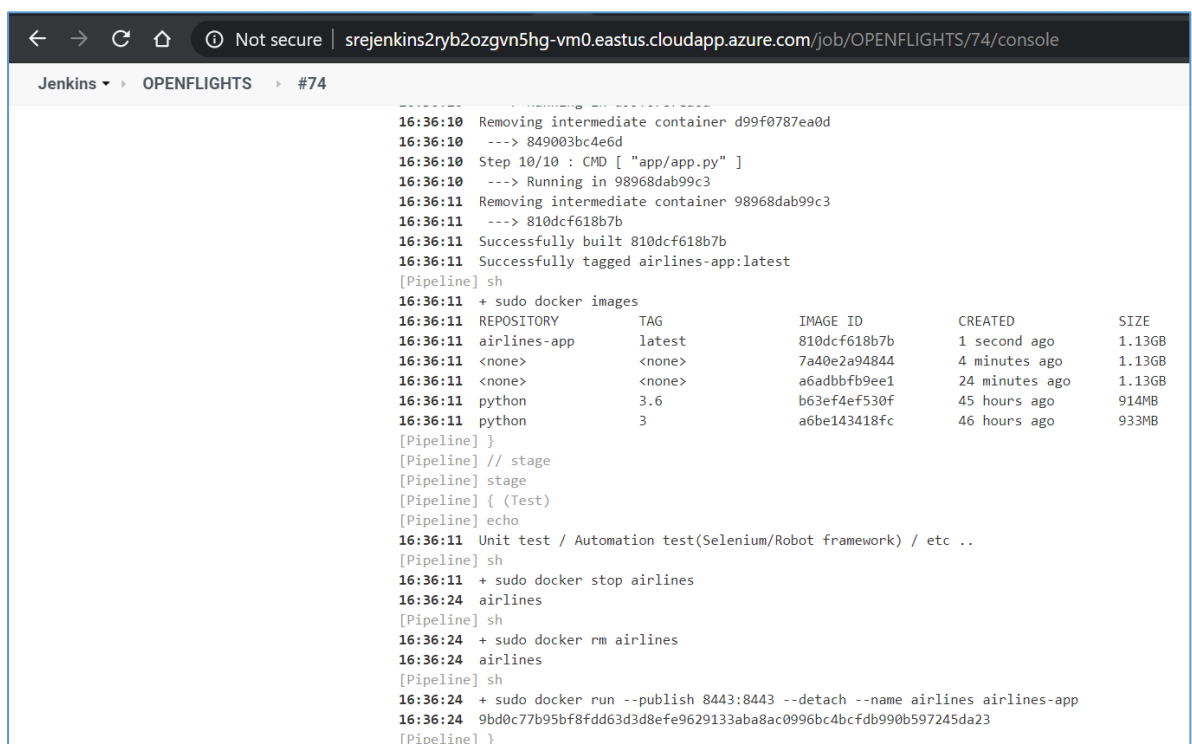
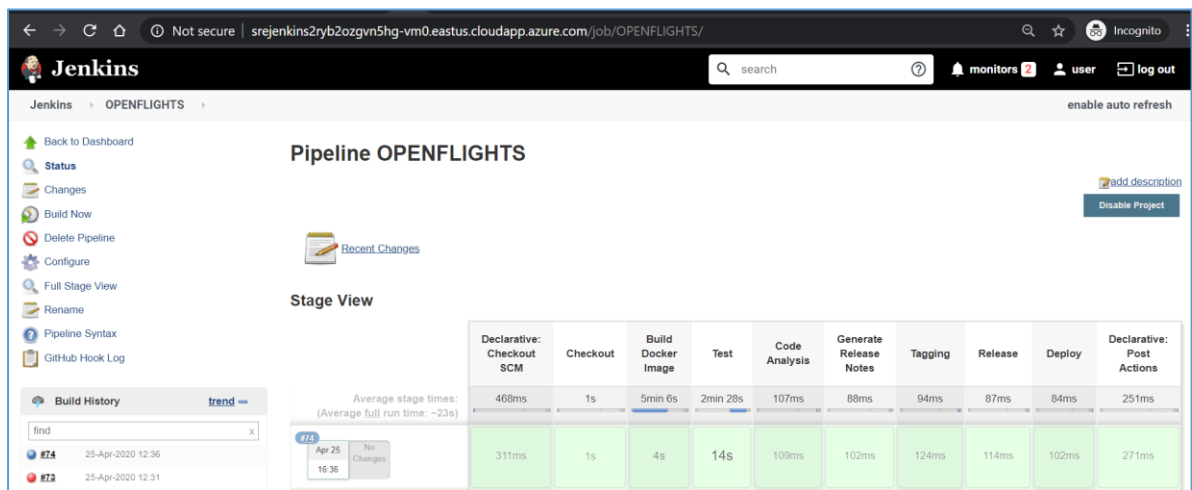


1. A GIT repo carrying the file parsing application code.
<https://github.com/nalubotula/openflightsproj.git>
2. A GIT repo carrying pipeline setup using pipeline as a code script (an added advantage).
<https://github.com/nalubotula/openflights.git>
3. Provisioned infrastructure should be well defined with proper tagging and naming convention.

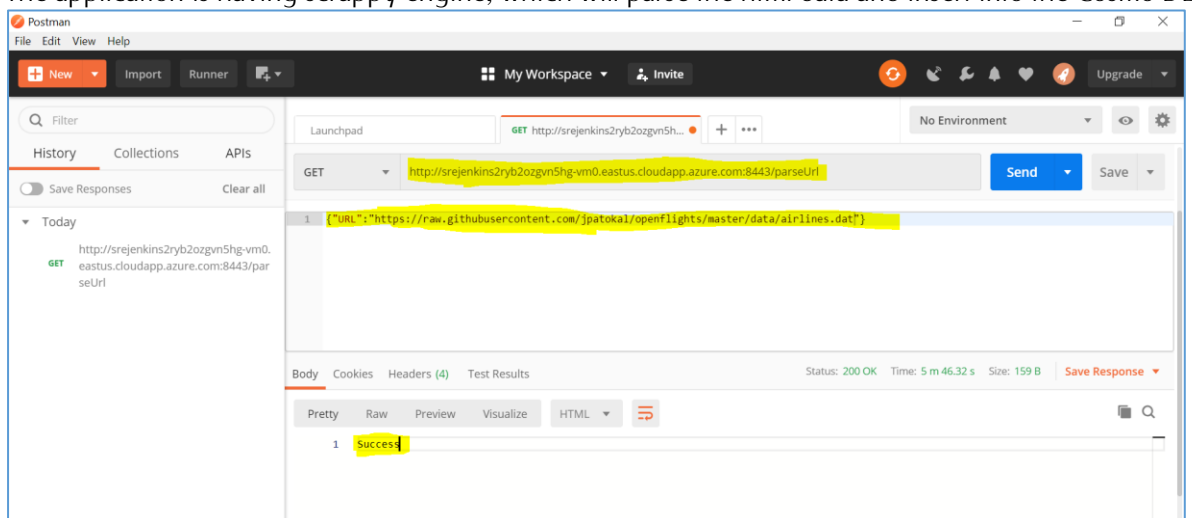
As we are using docker infrastructure, tagging and versioning is done during the build process which carries into the container registry

Infrastructure tags and naming conventions propagates form the code git repositories/branches to Jenkins as parameters, then to the hosting infrastructure code.

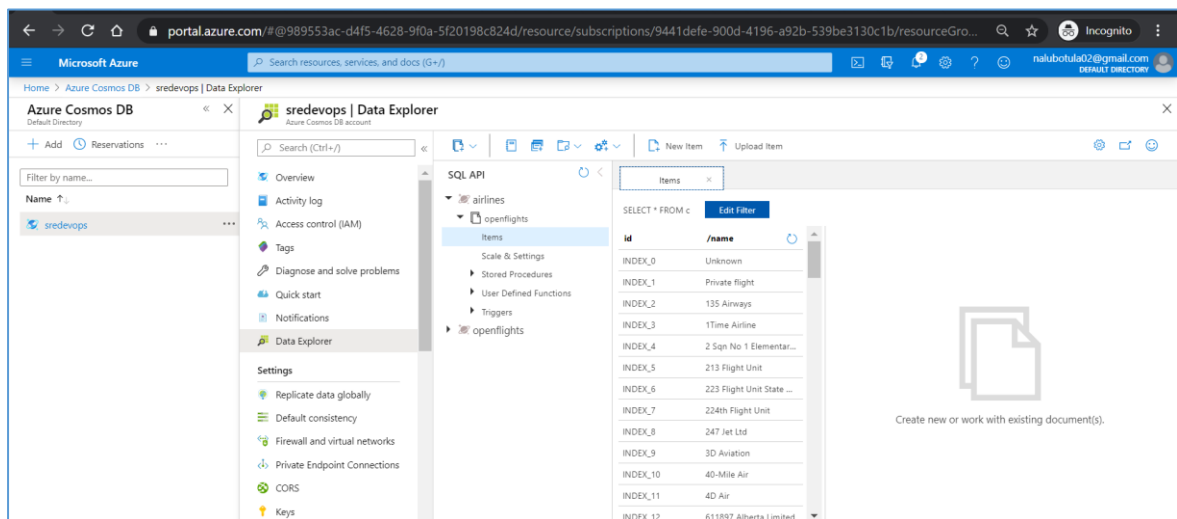
Application is auto built and deployed on a test server as a docker image [snapshots below]



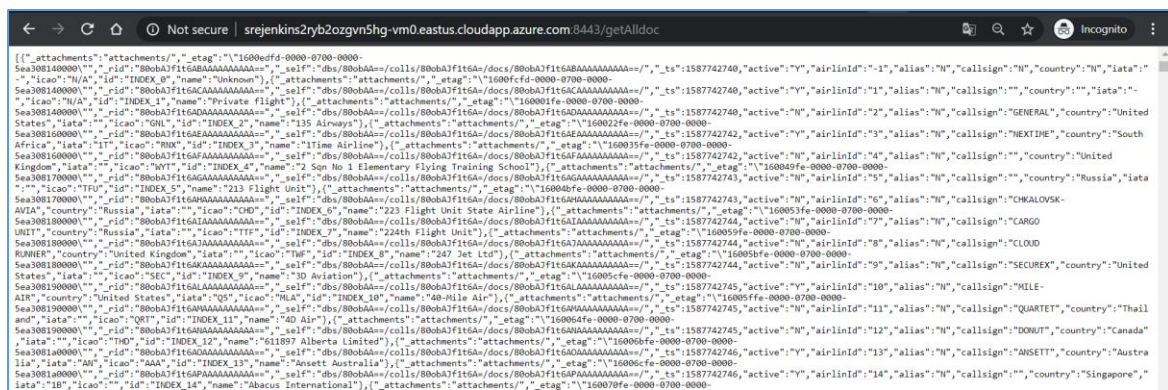
We will hit the deployed app end point from the postman , with the provide data file URL
 The application is having scrappy engine, which will parse the html data and insert into the Cosmo DB.



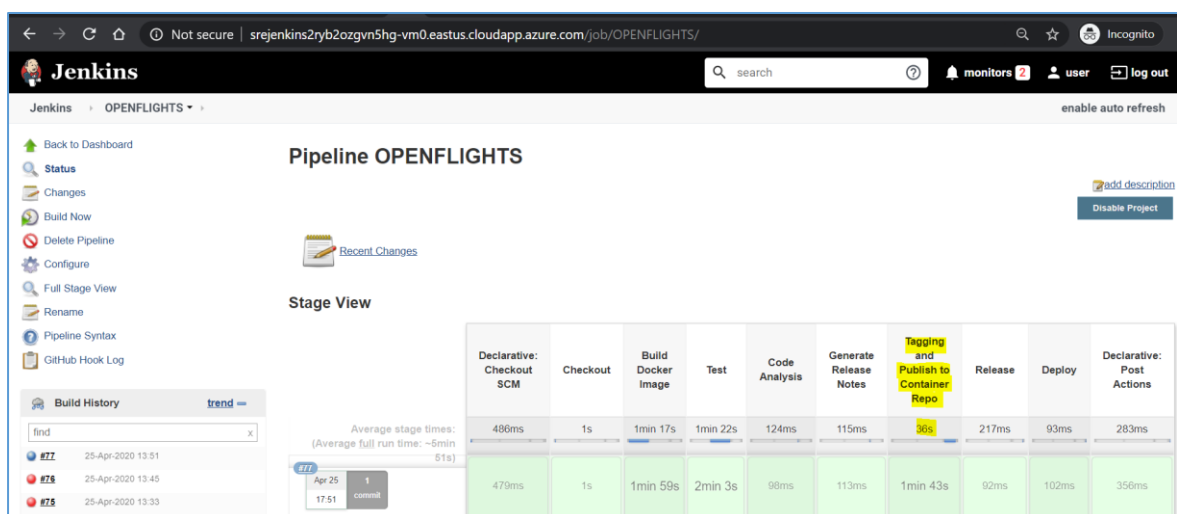
We can see the data been inserted into the Azure Cosmos DB [snapshots are provided below]



The application also has the service to fetch the data , which can be viewed from the browser on below mentioned URL.

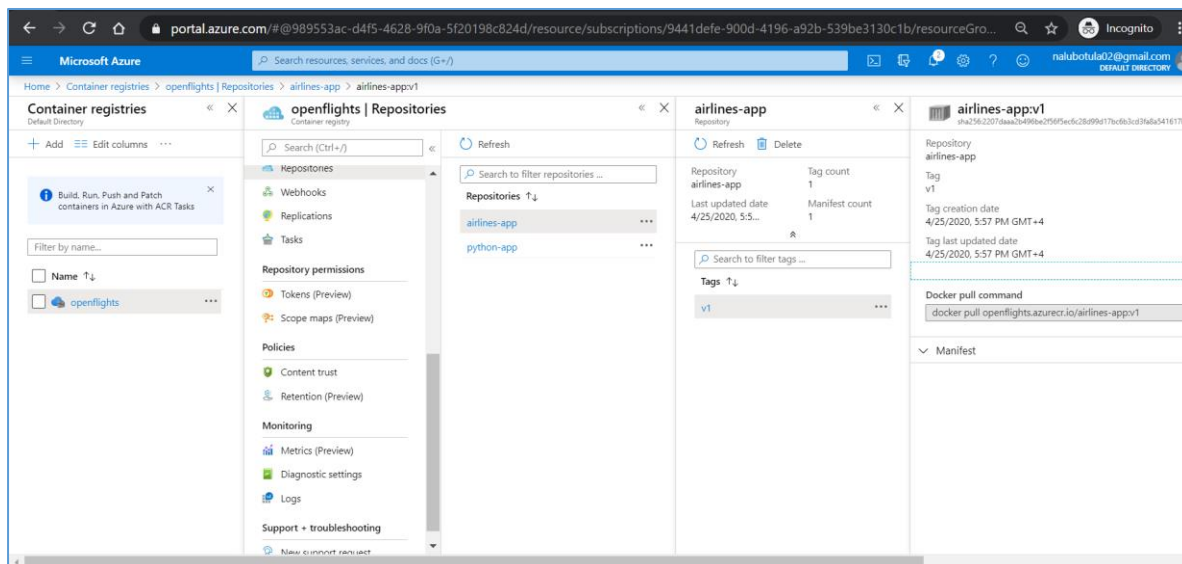


The tested application image is now tagged and pushed to Azure Container Registry [We can see in the snapshots]



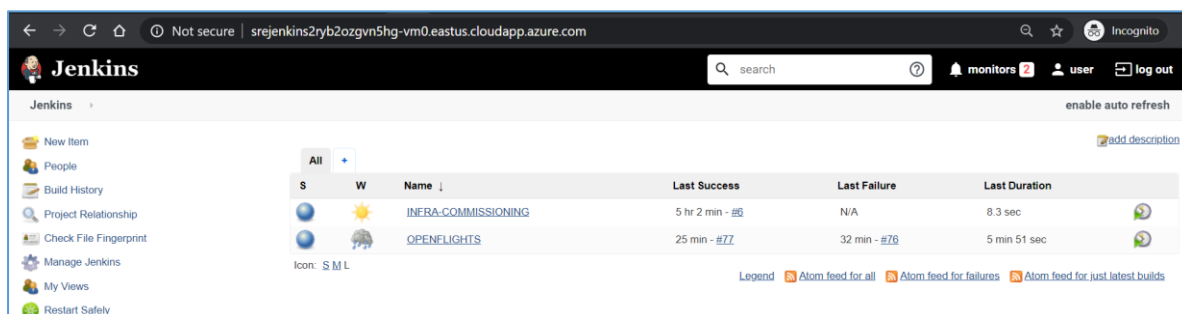
```
← → ↺ ↻ 🔒 Not secure | srejenkins2ryb2ozgvn5hg-vm0.eastus.cloudapp.azure.com/job/OPENFLIGHTS/77/console
Jenkins ▾ OPENFLIGHTS ▾ #77

[Pipeline] stage
[Pipeline] { (Tagging and Publish to Container Repo)
[Pipeline] echo
17:55:50 Tagging and Publish to Container Repo ..
[Pipeline] sh
17:55:50 + sudo docker rmi openflights.azurecr.io/airlines-app:v1
17:55:52 Untagged: openflights.azurecr.io/airlines-app:v1
17:55:52 Deleted: sha256:46ae7ed7e00c8808c0cddd64f917b5f6fe87e21c558546000de1000cf9add97f
17:55:52 Deleted: sha256:dc2ff20437cf5642960978793afe2eaa979834ca231fd7512192b80d60f0c56c
17:55:52 Deleted: sha256:5a16a427486e008cdfb654f2c54870b67e3d10ed3aa81d57b15f33efe39ac022
17:55:52 Deleted: sha256:481e8a43921480a24bbc93f809df7596a69bf79d383e82c6189a7c928d65b849
17:55:52 Deleted: sha256:e085eff1dda23b61c4589284b54ceef6e334fa96ce007ab02b933201a29b9e8
[Pipeline] sh
17:55:52 + sudo docker tag airlines-app openflights.azurecr.io/airlines-app:v1
[Pipeline] sh
17:55:57 + sudo az acr login --name openflights.azurecr.io
17:55:58 The login server endpoint suffix '.azurecr.io' is automatically omitted.
17:55:58 Login Succeeded
[Pipeline] sh
17:55:59 + sudo docker push openflights.azurecr.io/airlines-app:v1
17:55:59 The push refers to repository [openflights.azurecr.io/airlines-app]
17:55:59 97ddb8c6c44: Preparing
17:55:59 3ce232c5f6a2: Preparing
17:55:59 945616fb1a69: Preparing
17:55:59 6cd4fc6577f5: Preparing
17:56:11 6cd4fc6577f5: Pushed
17:56:17 b544d7bb9107: Pushed
17:56:25 baf481fca4b7: Pushed
17:56:51 9794a3b3ed45: Pushed
17:56:51 5f77a51ade6a: Pushed
17:56:51 8967306e673e: Pushed
17:57:09 3d3e92e98337: Pushed
17:57:27 e40d297cf5f8: Pushed
17:57:31 v1: digest: sha256:2207daaa2b496be2f56f5ec6c28d99d17bc6b3cd3fa8a541617b7724e3535eea size: 3061
```



We will further deploy this application into an Azure AKS Cluster

1. Infrastructure also will be provisioned via terraform scripts from a github repo.
2. Jenkins Job will initiate the provisioning.



Not secure | srejenkins2ryb2ozgvn5hg-vm0.eastus.cloudapp.azure.com/job/INFRA-COMMISSIONING/

Jenkins

search monitors 2 user log out

Jenkins ▾ INFRA-COMMISSIONING enable auto refresh

Back to Dashboard
Status
Changes
Build Now
Delete Pipeline
Configure
Full Stage View
Rename
Pipeline Syntax

Pipeline INFRA-COMMISSIONING

This Job will Commission the AKS Infrastructure from Pipeline Code

edit description
Disable Project

Recent Changes

Stage View

Build History trend

find

25-Apr-2020 09:14

Atom feed for all Atom feed for failures

Average stage times:
(Average full run time: ~8s)

Declarative: Checkout SCM	Checkout	Infrastructure De-commissioning	Infrastructure commissioning (AKS Cluster on Azure)	Declarative: Post Actions
525ms	1s	90ms	3s	290ms

Apr 25 13:14 1 commit

Not secure | srejenkins2ryb2ozgvn5hg-vm0.eastus.cloudapp.azure.com/job/INFRA-COMMISSIONING/6/console

Jenkins ▾ INFRA-COMMISSIONING #6

```
[Pipeline] stage
[Pipeline] { (Infrastructure commissioning (AKS Cluster on Azure))
[Pipeline] echo
13:14:41 Commissioning (AKS Cluster on Azure) ..
[Pipeline] sh
13:14:41 + terraform --version
13:14:41 Terraform v0.12.24
[Pipeline] sh
13:14:41 + cd flightservice/terraform-aks-k8s
13:14:41 + sudo terraform init -backend-config=storage_account_name=sredevops -backend-config=container_name=tfstate -backend-config=access_key=HaGZTn0LXMKLSIFduIf3+F8dzgDqW4Be0zHk337ac3mGVaZG02Rw5QZQd0AehhY6j67qkg7HE1wJQTa1BQ== -backend-config=key=codelab.microsoft.tfstate
13:14:41 [0m][1m]Initializing the backend...[0m
13:14:41 [0m][32m
13:14:41 Successfully configured the backend "azure"! Terraform will automatically
13:14:41 use this backend unless the backend configuration changes.[0m
13:14:41 [0m][1m]Initializing provider plugins...[0m
13:14:41 - Checking for available provider plugins...
13:14:42 - Downloading plugin for provider "azurerm" (hashicorp/azurerm) 2.7.0...
13:14:43 - Downloading plugin for provider "random" (hashicorp/random) 2.2.1...
13:14:44 The following providers do not have any version constraints in configuration,
13:14:44 so the latest version was installed.
13:14:44 To prevent automatic upgrades to new major versions that may contain breaking
13:14:44 changes, it is recommended to add version = "..." constraints to the
13:14:44 corresponding provider blocks in configuration, with the constraint strings
13:14:44 suggested below.
13:14:44 * provider.random: version = "~> 2.2"
13:14:44 [0m][1m][32mTerraform has been successfully initialized!
13:14:44 [0m][32m[0m
```

We can see the kubernetes cluster created in the Azure as AKS Service.

Current node size is 2

portal.azure.com/#@989553ac-d4f5-4628-9f0a-5f20198c824d/resource/subscriptions/9441defe-900d-4196-a92b-539be3130c1b/resourceGro...

Microsoft Azure Search resources, services, and docs (G+)

Home > Resource groups > k8sredevops

Resource groups

+ Add Manage view

Filter by name...

Name ↑

- akscontainer
- cloud-shell-storage-centralindia
- DevOps
- k8sredevops
- MC_k8sredevops_k8sredevops_eastus2
- NetworkWatcherRG
- srejenkins

Page 1 of 1

k8sredevops

Overview

Activity log

Access control (IAM)

Tags

Events

Settings

Quickstart

Deployments

Policies

Properties

Locks

Export template

Cost Management

Cost analysis

Cost alerts (preview)

Budgets

Subscription (change): Azure subscription 1

Subscription ID : 9441defe-900d-4196-a92b-539be3130c1b

Tags (change) : Click here to add tags

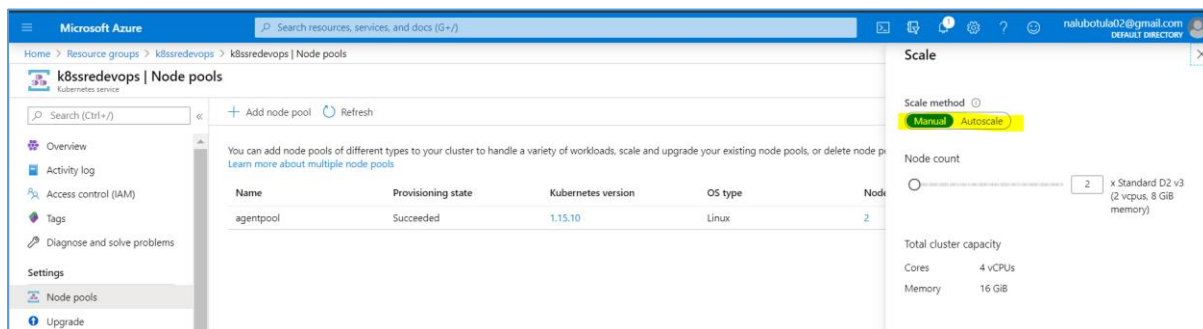
Filter by name... Type == all Location == all Add filter

Showing 1 to 3 of 3 records. Show hidden types

Name ↑	Type ↑	Location ↑
ContainerInsights(testloganalyticsworkspaceName-2510717841218586885)	Solution	East US
k8sredevops	Kubernetes service	East US 2
testLogAnalyticsWorkspaceName-2510717841218586885	Log Analytics workspace	East US

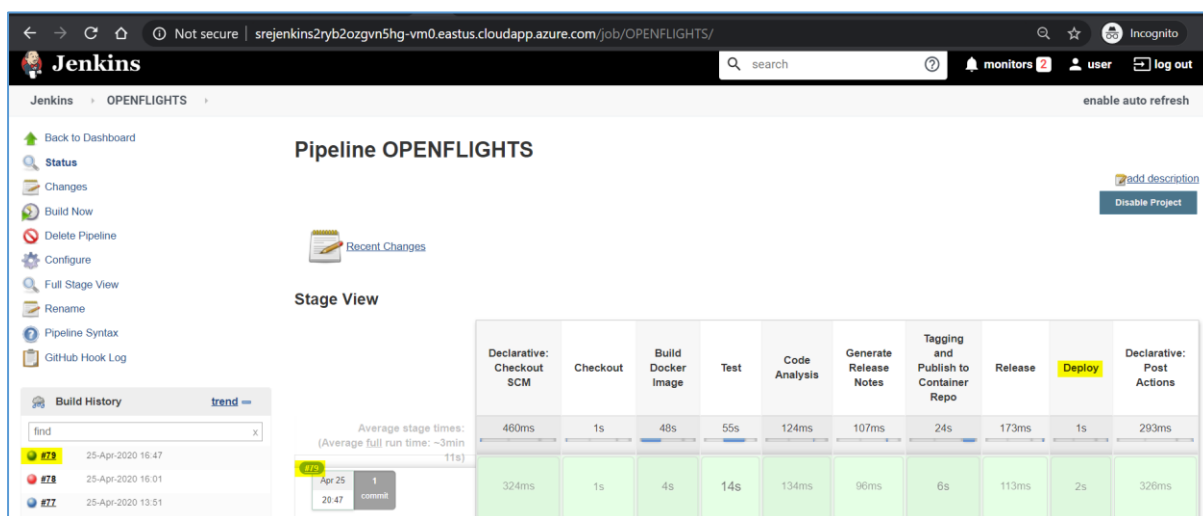
< Previous Page 1 of 1 Next >

We can upscale the nodes based on load, or enable auto-scaling option.



Now I have deployed the application to AKE Cluster , using deploy.yaml [code is available in the repo]

Configured pipeline script, to execute as part of Jenkins pipeline.



```

[Pipeline] stage
[Pipeline] { (Deploy)
[Pipeline] echo
20:48:25 Deploy the docker image from ACR to AKE CLuster ..
[Pipeline] sh
20:48:25 + sudo kubectl get nodes
20:48:27 NAME STATUS ROLES AGE VERSION
20:48:27 aks-agentpool-18820554-vmss000000 Ready agent 135m v1.15.10
20:48:27 aks-agentpool-18820554-vmss000001 Ready agent 135m v1.15.10
[Pipeline] sh
20:48:27 + sudo kubectl apply -f deploy.yaml
20:48:27 service/airlines-app created
20:48:27 deployment.apps/airlines-front created
20:48:27 service/airlines-front created
[Pipeline] }
[Pipeline] // stage
[Pipeline] stage
[Pipeline] { (Declarative: Post Actions)
[Pipeline] script
[Pipeline] {
[Pipeline] }
[Pipeline] // script
[Pipeline] echo
20:48:28 Status Changed: [From: FAILURE, To: SUCCESS]
[Pipeline] echo
20:48:28 SUCCESS
[Pipeline] }
[Pipeline] // stage
[Pipeline] }
[Pipeline] // timestamps
[Pipeline] }
[Pipeline] // timeout
[Pipeline] }
[Pipeline] // withEnv
[Pipeline] }

```

Since I have not installed any kubernetes monitoring tools, I had to check backend.


```

root@srejenkins0:/home/bitnami# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
airlines-front-7d75d7bd6c-ttxcx    0/1     ContainerCreating   0           6m46s
root@srejenkins0:/home/bitnami# kubectl get events
LAST SEEN   TYPE      REASON              OBJECT                                          MESSAGE
7m49s       Normal    Scheduled            pod/airlines-front-7d75d7bd6c-ttxcx           Successfully assigned default/airlines-front-7d75d7bd6c-ttxcx to aks-agentpool-18820554-vmss000001
7m47s       Normal    Pulling              pod/airlines-front-7d75d7bd6c-ttxcx           Pulling image "openflights.azurecr.io/airlines-app:v1"
63s         Normal    Pulled               pod/airlines-front-7d75d7bd6c-ttxcx           Successfully pulled image "openflights.azurecr.io/airlines-app:v1"
61s         Normal    Created              pod/airlines-front-7d75d7bd6c-ttxcx           Created container airlines-app
61s         Normal    Started              pod/airlines-front-7d75d7bd6c-ttxcx           Started container airlines-app
7m49s       Normal    SuccessfulCreate     replicaset/airlines-front-7d75d7bd6c-ttxcx    Created pod: airlines-front-7d75d7bd6c-ttxcx
7m49s       Normal    ScalingReplicaSet    deployment/airlines-front                      Scaled up replica set airlines-front-7d75d7bd6c to 1
7m49s       Normal    EnsuringLoadBalancer service/airlines-front                         Ensuring load balancer
7m6s        Normal    EnsuredLoadBalancer service/airlines-front                         Ensured load balancer
root@srejenkins0:/home/bitnami# kubectl get logs
error: the server doesn't have a resource type "logs"
root@srejenkins0:/home/bitnami# tail /var/log/pods/*
tail: cannot open '/var/log/pods/*' for reading: No such file or directory
root@srejenkins0:/home/bitnami# kubectl logs airlines-app --tail 10
Error from server (NotFound): pods "airlines-app" not found
root@srejenkins0:/home/bitnami# kubectl logs airlines-front-7d75d7bd6c-ttxcx --tail 10
* Serving Flask app "app" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://0.0.0.0:8443/ (Press CTRL+C to quit)
root@srejenkins0:/home/bitnami# kubectl get svc

```

```

root@srejenkins0:/home/bitnami# kubectl get svc
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
airlines-app        ClusterIP   10.0.245.107  <none>         6379/TCP         16m
airlines-front      LoadBalancer 10.0.37.69    40.70.86.203   80:31068/TCP     16m
kubernetes          ClusterIP   10.0.0.1      <none>         443/TCP          156m
root@srejenkins0:/home/bitnami# curl wgetip.com
40.71.170.195root@srejenkins0:/home/bitnami# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE                                NOMINATED NODE   READINESS GATES
airlines-front-7d75d7bd6c-ttxcx    1/1     Running   0           20m   10.244.0.5      aks-agentpool-18820554-vmss000001 <none>           <none>
root@srejenkins0:/home/bitnami#

```

Note: A WebHook has been created on the application git [github] repo, which initiates the Jenkins Build and Deploy process, on every commit into master branch

nalubotula / openflightsproj

Unwatch

1

Star

0

Fork

0

<> Code

Issues 0

Pull requests 0

Actions

Projects 0

Wiki

Security 0

Insights

Settings

Options

Manage access

Branches

Webhooks

Webhooks

Add webhook

Webhooks allow external services to be notified when certain events happen. When the specified events happen, we'll send a POST request to each of the URLs you provide. Learn more in our [Webhooks Guide](#).

✓

http://srejenkins2ryb2ozgvn5hg-vm0.eastus.cloudapp.azure.com/github-webhook/ (push)

Edit

Delete